

Linking

2017 – 2018

Florent Gluck – Florent.Gluck@hesge.ch

Version 0.3

Example

main.c

```
int sum(int *a, int n);

int array[2] = {1, 2};

int main() {
    int val = sum(array, 2);
    return val;
}
```

sum.c

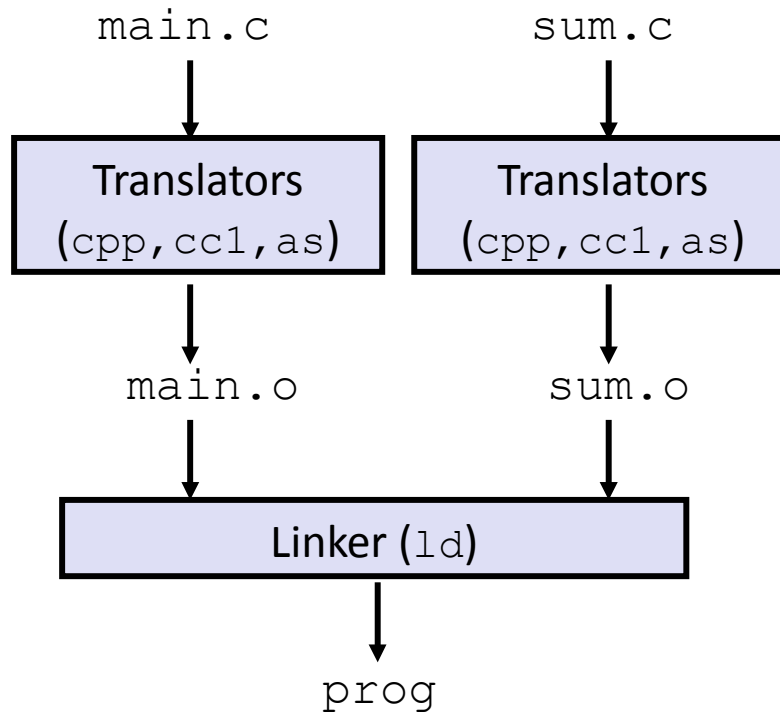
```
int sum(int *a, int n) {
    int i, s = 0;

    for (i = 0; i < n; i++) {
        s += a[i];
    }
    return s;
}
```

Linking

Programs are translated and linked using a toolchain:

```
gcc -o prog main.c sum.c  
./prog
```



Sources files

Relocatable object files

Fully linked **executable** object file
(contains code and data for all
functions defined in `main.c` and
`sum.c`)

Why linkers? (1)

Reason 1: **Modularity**

- Program can be written as a collection of smaller source files, rather than one monolithic mass.
- Can build libraries of common functions:
 - e.g., math library, standard C library

Why linkers? (2)

Reason 2 : **Efficiency**

- **Time: Separate compilation**
 - Change one source file, compile, and then relink.
 - No need to recompile other source files.
- **Space: libraries**
 - Common functions can be aggregated into a single file...
 - Yet executable files and running memory images contain only code for the functions they actually use.

What do linkers do? (1)

Step 1: symbol resolution

- Programs define and reference symbols (global variables and functions):

```
int sum(int *a, int n) {...} // define symbol sum
sum(...);                  // reference symbol sum
int *xp = &x;               // define symbol xp, reference x
```

- Symbol definitions are stored in the object file (by assembler) in the **symbol table**.
- Each entry includes name, size, and location of symbol.
- During **symbol resolution** step, the linker associates each symbol reference with exactly one symbol definition.

Three kinds of object files

- **Relocatable** object file (.o file)
 - Contains code and data in a form that can be combined with other **relocatable** object files to form an **executable** object file.
 - Each .o file is produced from exactly one source (.c) file
- **Executable** object file (a.out file)
 - Contains code and data in a form that can be copied directly into memory and then executed.
- **Shared** object file (.so file)
 - Special type of **relocatable** object file that can be loaded into memory and linked dynamically, at either load time or run-time.
 - Called Dynamic Link Libraries (DLLs) by Windows

Three kinds of object files

compiler and assembler

- **Relocatable** object file (.o file)
 - Contains code and data in a form that can be combined with other **relocatable** object files to form an **executable** object file.
 - Each .o file is produced from exactly one source (.c) file
- **Executable** object file (a.out file)
 - Contains code and data in a form that can be copied directly into memory and then executed.
- **Shared** object file (.so file)
 - Special type of **relocatable** object file that can be loaded into memory and linked dynamically, at either load time or run-time.
 - Called Dynamic Link Libraries (DLLs) by Windows

Three kinds of object files

- **Relocatable** object file (.o file)

- Contains code and data in a form that can be combined with other **relocatable** object files to form an **executable** object file.
- Each .o file is produced from exactly one source (.c) file

linker (ld)

- **Executable** object file (a.out file)

- Contains code and data in a form that can be copied directly into memory and then executed.

- **Shared** object file (.so file)

- Special type of **relocatable** object file that can be loaded into memory and linked dynamically, at either load time or run-time.
- Called Dynamic Link Libraries (DLLs) by Windows

Executable and Linkable Format (ELF)

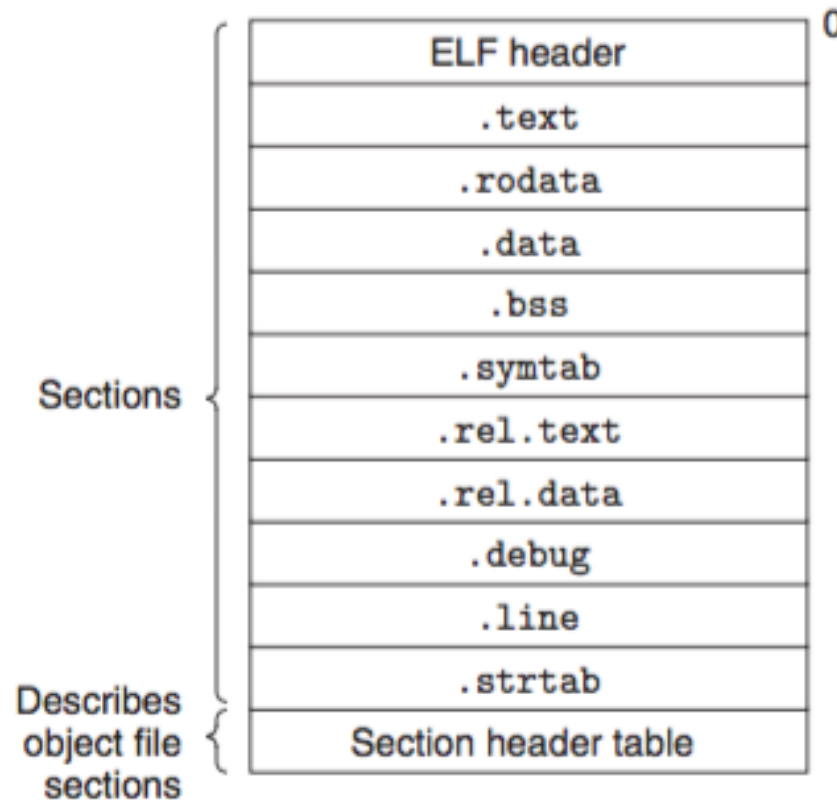
- ELF is an unified, architecture agnostic, binary file format designed to store object files content.
- Quick adoption by numerous operating systems and architectures.
- Derived from UNIX AT&T System V and chosen in 1999 to be the standard binary file format for all UNIX systems.
- One **unified** format for
 - **Relocatable** object files (.o),
 - **Executable** object files (a.out)
 - **Shared** object files (.so)
- Generic name: ELF binaries

Relocatable object files (1)

- A **relocatable** object file is basically a collection of sections.
- Each section contains a single type of information, such as: program code, read only data, symbols, etc.
- Each symbol's address is defined relatively to its section
- A function's entry point is relative to the section where the function's code resides.

Relocatable object files (2)

Example of a **relocatable** object file:

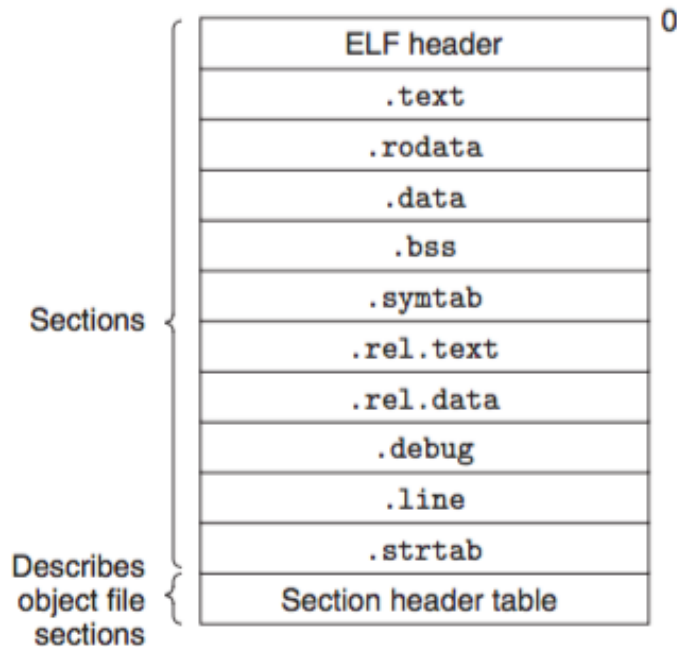


Executable object files

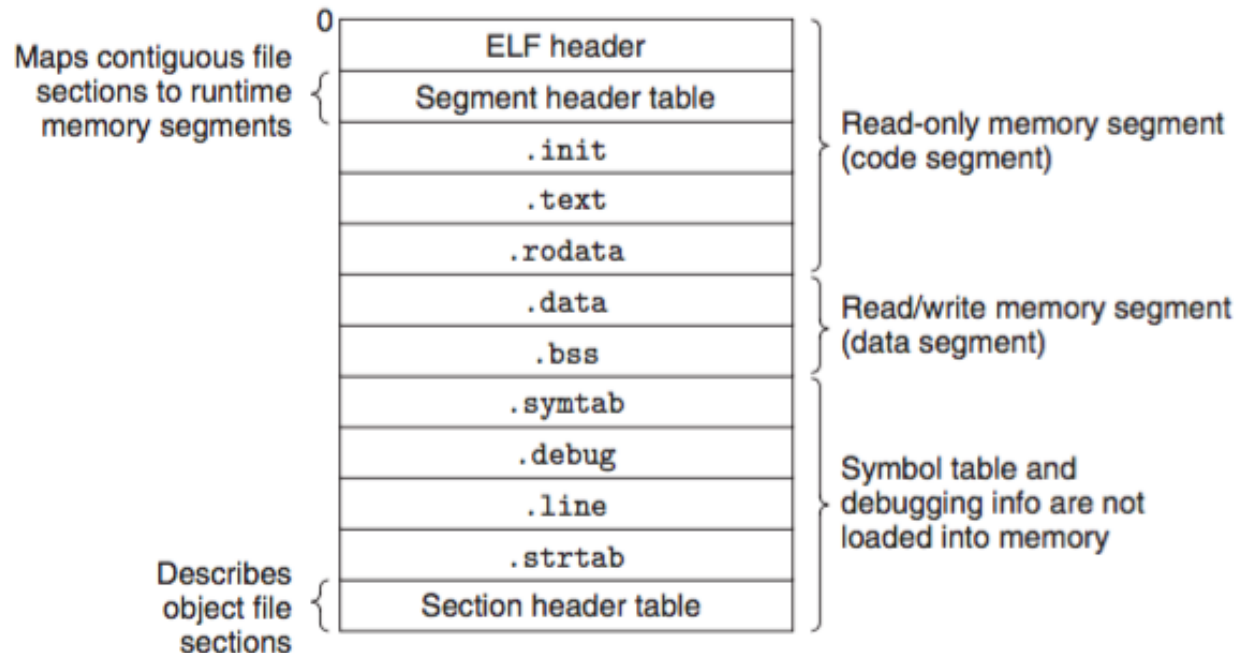
- **Very similar** to **relocatable** object files, except :
 - Program's entrypoint stored in the ELF header.
 - .text, .rodata, and .data sections **similar** to **relocatable** object files, except:
 - These sections' addresses are relocated at runtime.
- ELF format designed to easily load into RAM:
 - Contiguous chunks of files mapped to contiguous memory segments.
- The process of copying a program into RAM to be executed is called "loading".
 - Performed by the "loader" of the operating system.

Relocatable vs executable files

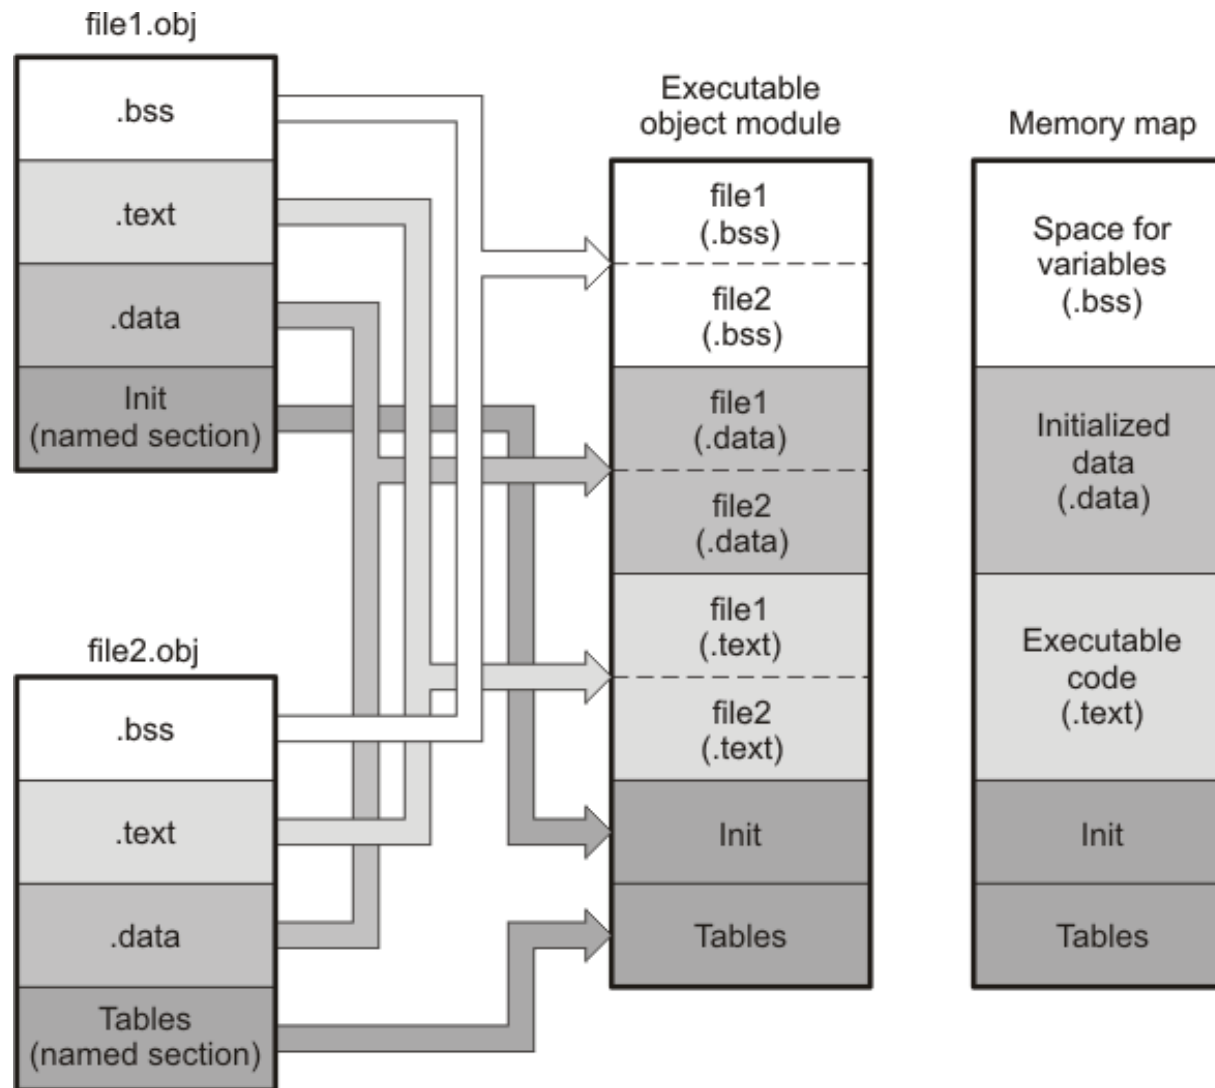
Relocatable object file:



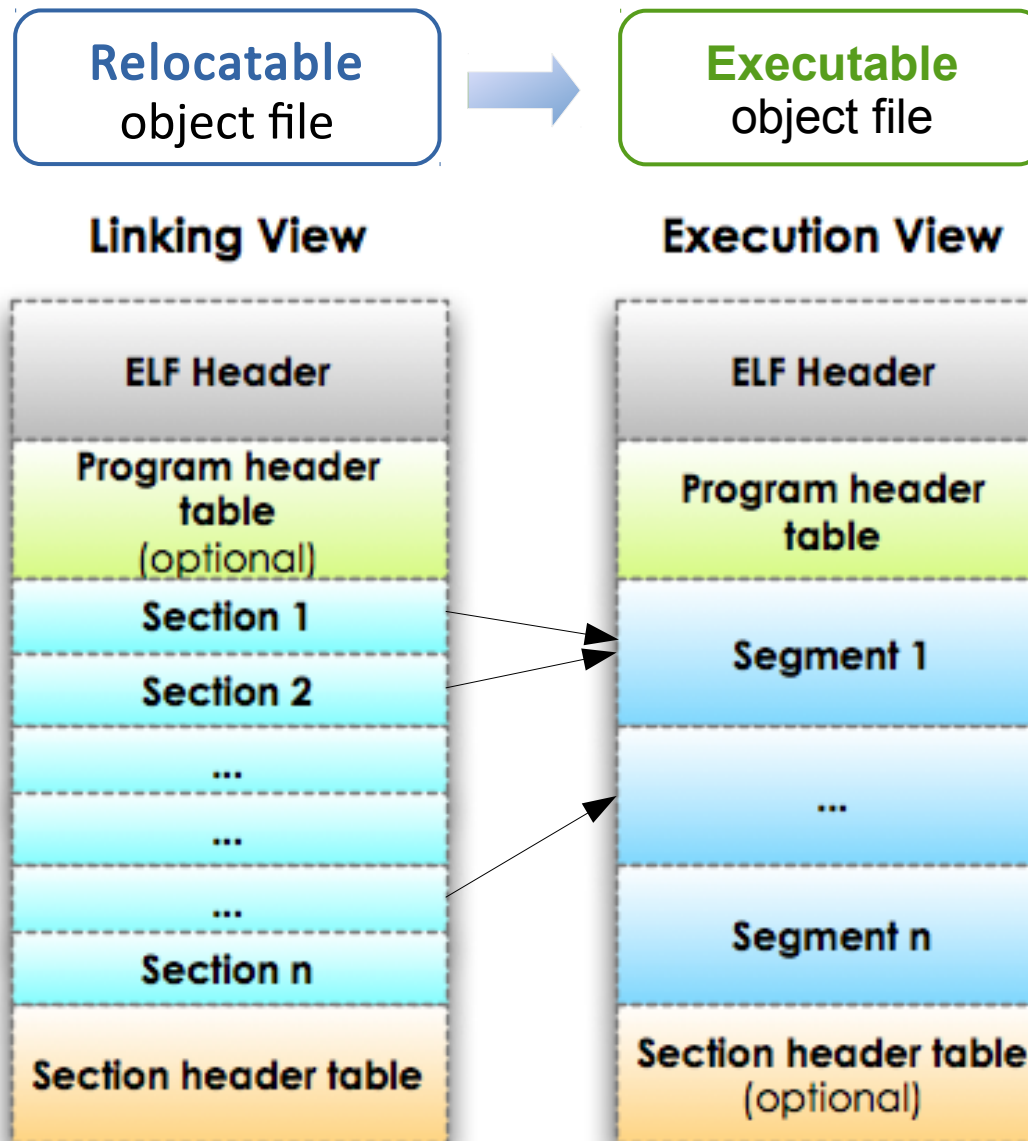
Executable object file:



From relocatable to executable (1)



From relocatable to executable (2)



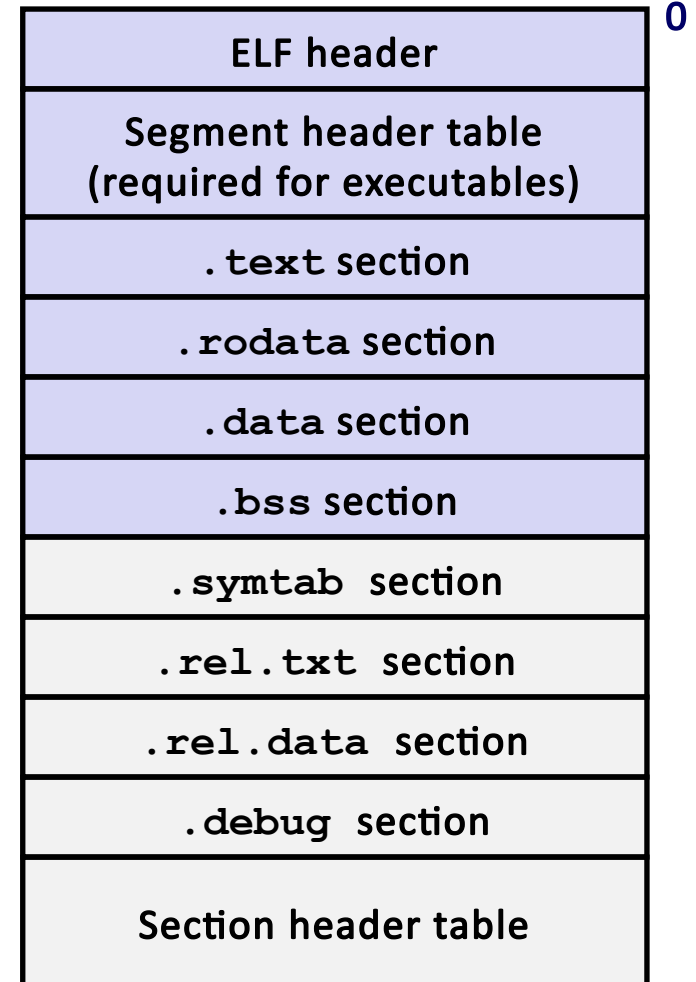
Structure of an ELF file

An ELF file has two personalities:

- **A link view**: compilers, assembler and linkers treat an ELF file as a set of logical **sections** defined by a **section header table**.
- **An execution view**: the loader treats an ELF file as a set of **segments** defined by a **program header table**.

ELF file format (1)

- **Elf header**
 - word size, byte ordering, file type (.o, exec, .so), machine type, etc.
- **Segment header table**
 - page size, virtual addresses memory segments (sections), segment sizes
- **.text section**
 - code
- **.rodata section**
 - read only data
- **.data section**
 - initialized global variables
- **.bss section**
 - uninitialized global variables
 - “Block Started by Symbol”
 - “Better Save Space”
 - has section header but occupies no space



ELF file format (2)

- **.symtab section**
 - symbol table
 - procedure and static variable names
 - section names and locations
- **.rel.text section**
 - relocation info for .text section
 - addresses of instructions that will need to be modified in the executable
 - instructions for modifying.
- **.rel.data section**
 - relocation info for .data section
 - addresses of pointer data that will need to be modified in the merged executable
- **.debug section**
 - info for symbolic debugging (gcc -g)
- **Section header table**
 - offsets and sizes of each section

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.text section
.rel.data section
.debug section
Section header table

0

Linker symbols

- **Global symbols**

- Symbols defined by module **m** that can be referenced by other modules.
- E.g.: `non-static` C functions and `non-static` global variables.

- **External symbols**

- Global symbols that are referenced by module **m** but defined by some other module.

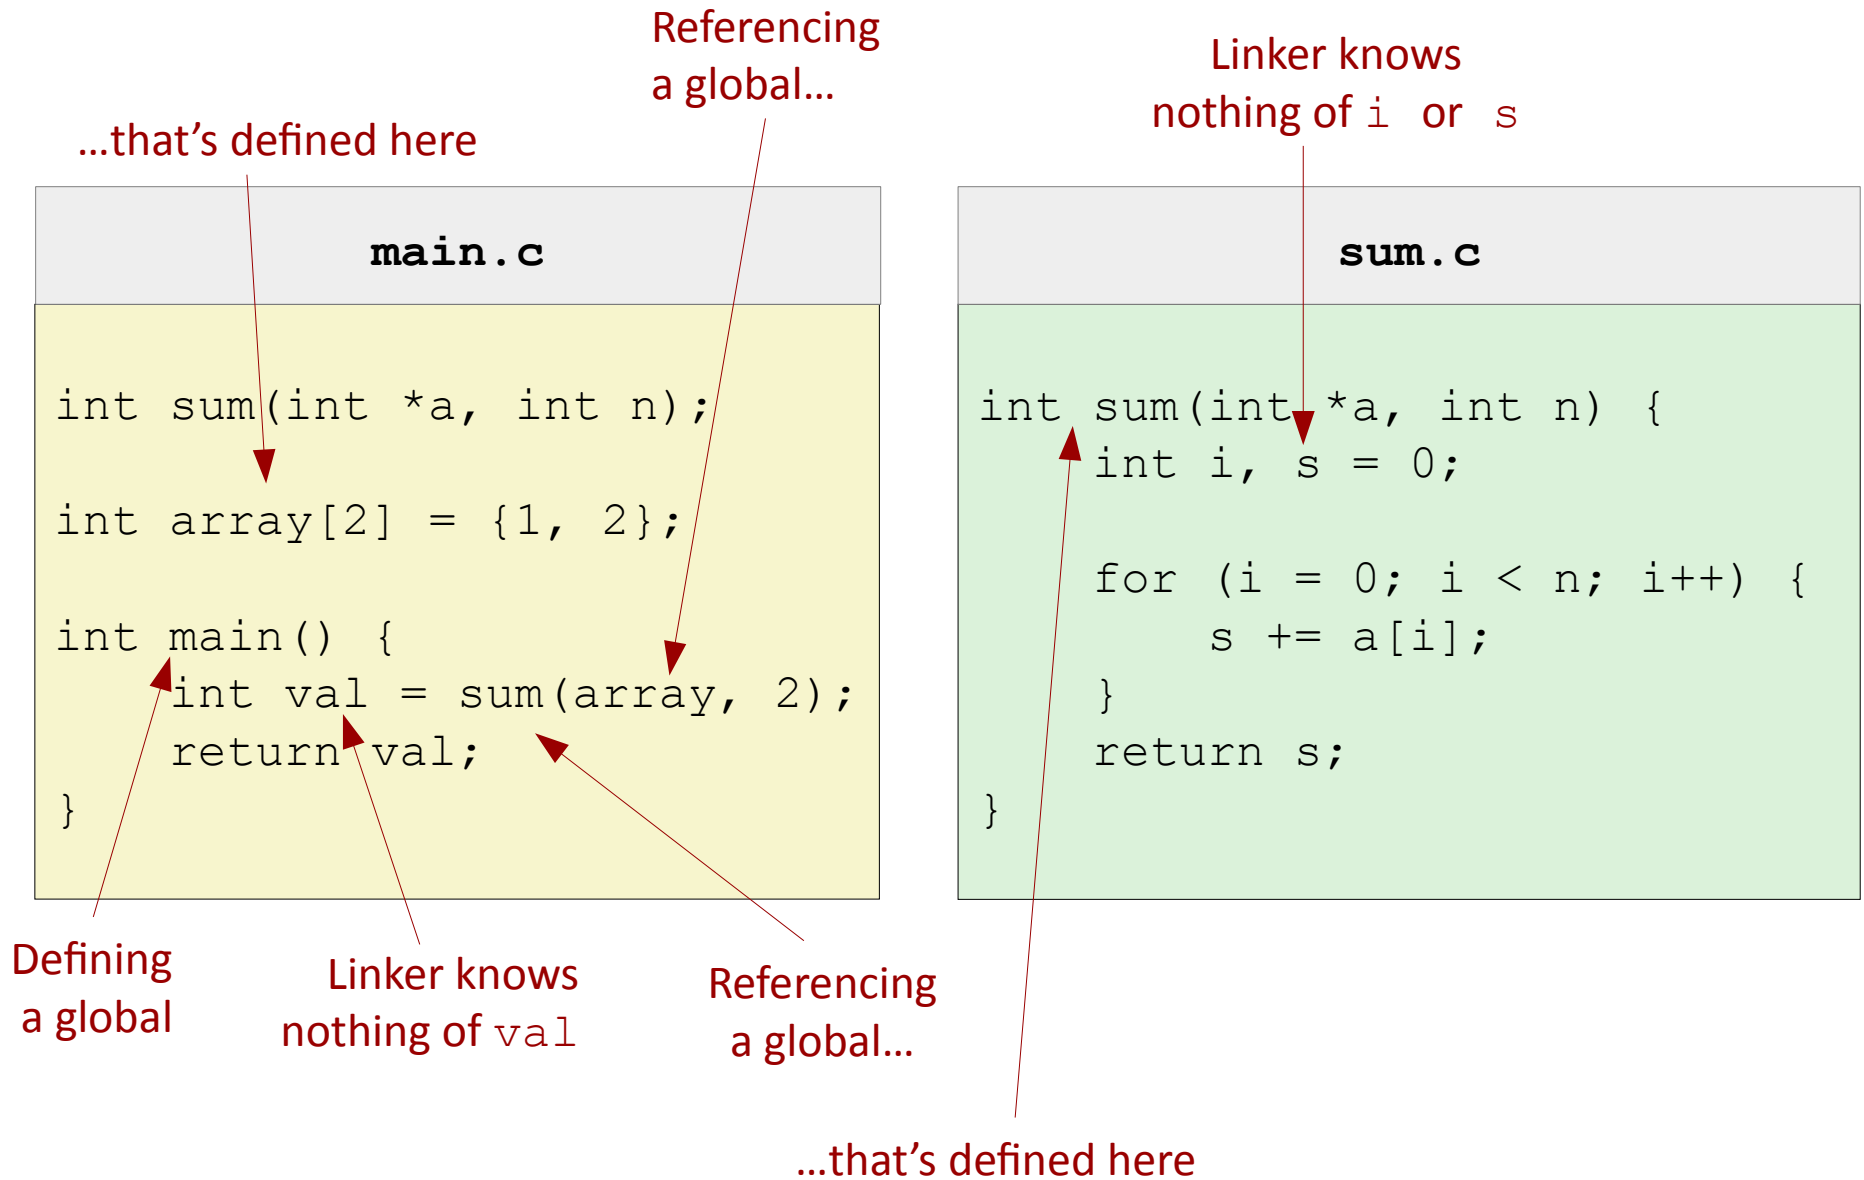
- **Local symbols**

- Symbols that are defined and referenced exclusively by module **m**.
- E.g.: C functions and global variables defined with the `static` attribute.
- Local linker symbols are not local program variables

Executable binary generation and execution

- Two steps are performed when generating an executable binary file:
 - Step 1: Symbol resolution
 - Step 2: Relocation
- When the binary file is executed, the **loader** loads it in memory (RAM)

Step 1: symbol resolution



Local symbols



- Local non-static C variables vs. local static C variables
 - local non-static C variables: stored on the stack
 - local static C variables: stored in either .bss, or .data

```
int f() {  
    static int x = 0;  
    int y = 3;  
    return x+y;  
}  
  
int g() {  
    static int x = 1;  
    return x;  
}
```



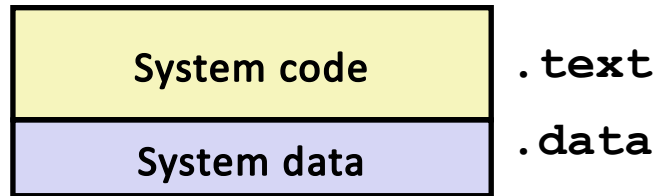
- Compiler allocates space in .data for each definition of x
- Creates local symbols in the symbol table with unique names, e.g., x.1 and x.2.

Global variables

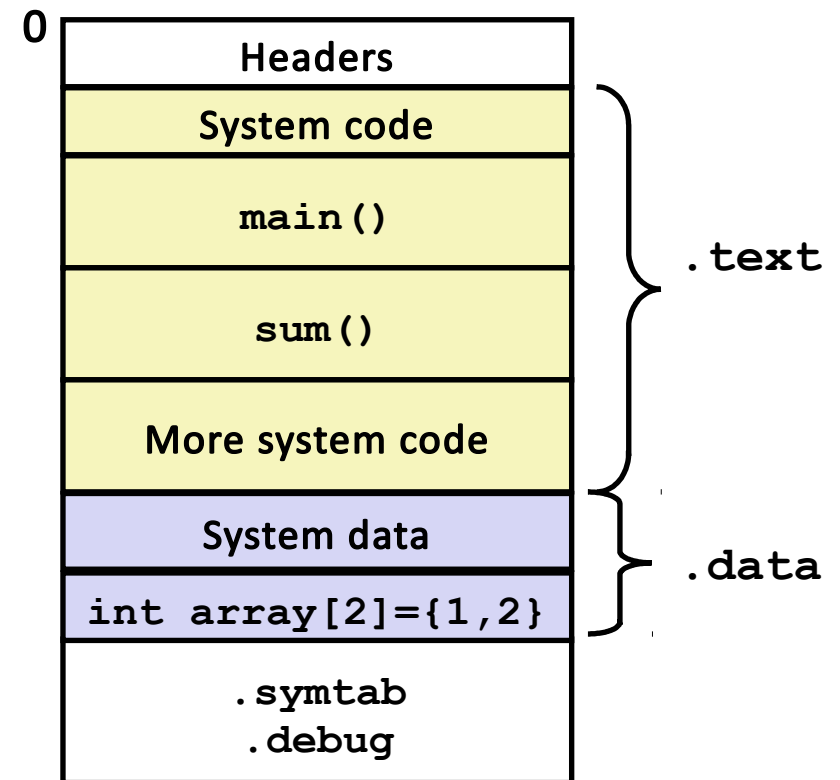
- **Avoid** if you can!
- Otherwise
 - Use `static` if you can
 - Initialize if you define a global variable
 - Use `extern` if you reference an external global variable

Step 2: relocation

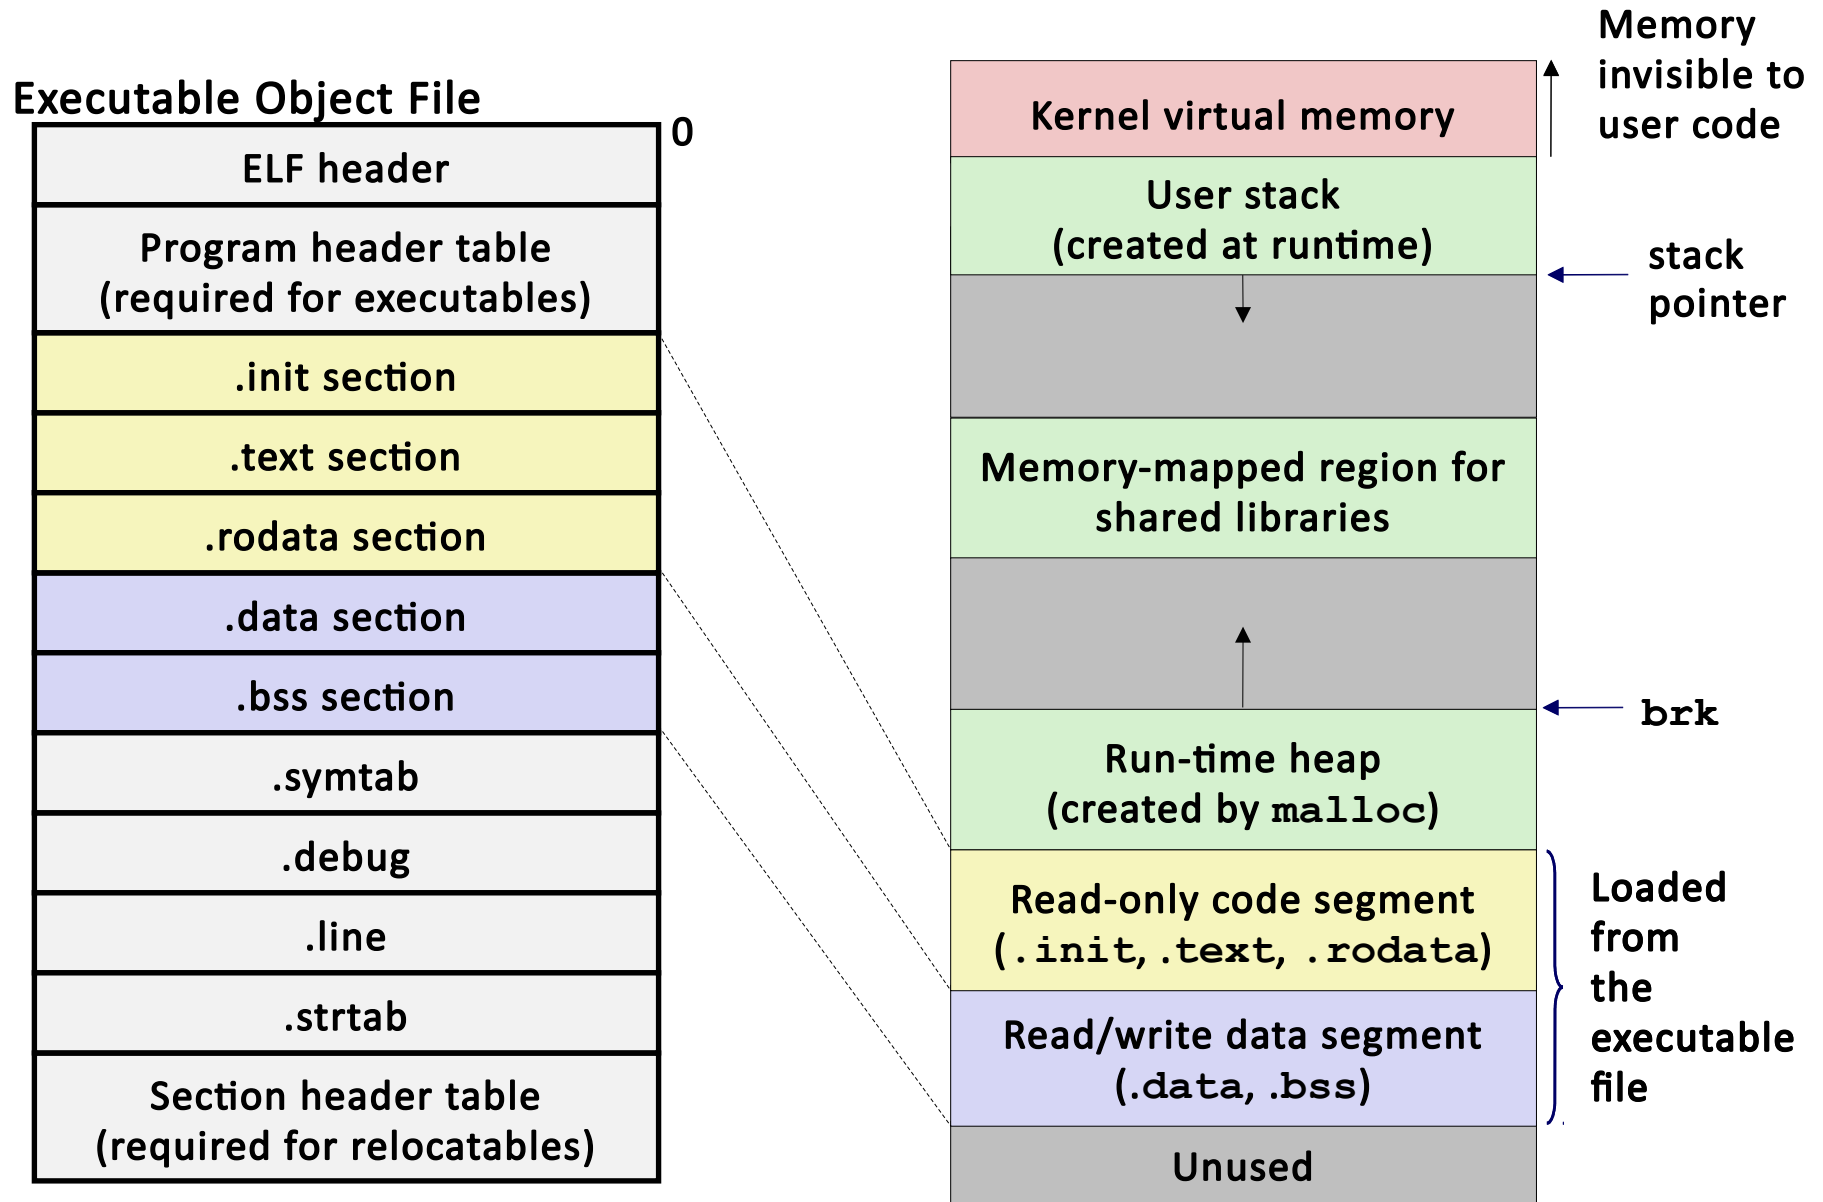
Relocatable Object Files



Executable Object File



Loading executable object files



readelf tool

- To inspect the content of ELF files, use the `readelf` tool.
- Exemple with the `test.c` file below:

```
int A;  
int B = 10;  
static int C = 4;  
  
int pipo() { return 42; }  
  
void main() { printf("blahblah\n"); }
```

ELF header – relocatable file

```
$ gcc -c test.c
```

```
$ readelf -h test.o
```

ELF Header:

```
Magic: 7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
Class: ELF64
Data: 2's complement, little endian
Version: 1 (current)
OS/ABI: UNIX - System V
ABI Version: 0
Type: REL (Relocatable file)
Machine: Advanced Micro Devices X86-64
Version: 0x1
Entry point address: 0x0
Start of program headers: 0 (bytes into file)
Start of section headers: 848 (bytes into file)
Flags: 0x0
Size of this header: 64 (bytes)
Size of program headers: 0 (bytes)
Number of program headers: 0
Size of section headers: 64 (bytes)
Number of section headers: 13
Section header string table index: 10
```

ELF header – **executable** file

```
$ gcc test.c -o test
$ readelf -h test
```

```
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                               2's complement, little endian
  Version:                             1 (current)
  OS/ABI:                               UNIX - System V
  ABI Version:                           0
  Type:                                EXEC (Executable file)
  Machine:                             Advanced Micro Devices X86-64
  Version:                               0x1
  Entry point address:                0x400430
  Start of program headers:              64 (bytes into file)
  Start of section headers:              6728 (bytes into file)
  Flags:                                0x0
  Size of this header:                    64 (bytes)
  Size of program headers:                56 (bytes)
  Number of program headers:               9
  Size of section headers:                64 (bytes)
  Number of section headers:              31
  Section header string table index:     28
```

Symbole table

```
$ readelf -s test.o
```

Symbol table '.symtab' contains 15 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	0000000000000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	0000000000000000	0	FILE	LOCAL	DEFAULT	ABS	test.c
2:	0000000000000000	0	SECTION	LOCAL	DEFAULT	1	
3:	0000000000000000	0	SECTION	LOCAL	DEFAULT	3	
4:	0000000000000000	0	SECTION	LOCAL	DEFAULT	4	
5:	0000000000000004	4	OBJECT	LOCAL	DEFAULT	3	C
6:	0000000000000000	0	SECTION	LOCAL	DEFAULT	5	
7:	0000000000000000	0	SECTION	LOCAL	DEFAULT	7	
8:	0000000000000000	0	SECTION	LOCAL	DEFAULT	8	
9:	0000000000000000	0	SECTION	LOCAL	DEFAULT	6	
10:	0000000000000004	4	OBJECT	GLOBAL	DEFAULT	COM	A
11:	0000000000000000	4	OBJECT	GLOBAL	DEFAULT	3	B
12:	0000000000000000	11	FUNC	GLOBAL	DEFAULT	1	pipo
13:	000000000000000b	17	FUNC	GLOBAL	DEFAULT	1	main
14:	0000000000000000	0	NOTYPE	GLOBAL	DEFAULT	UND	puts

- ABS specifies non relocatable symbols.
- UND specifies non resolved symbols (i.e. defined elsewhere).
- COM specified a non allocated object.

Tools for inspecting ELF files

- `nm` to list symbols from object files
- `readelf` to displays information about ELF files
- `objdump` to display information from object files
- To display the usage:
 - `nm --help`
 - `readelf --help`
 - `objdump --help`

Static vs dynamic linking

- **Static** linking is the result of the linker copying all library routines used in the program into the executable image.
 - Require more disk space and memory than dynamic linking
 - Faster and more portable, since it does not require the presence of the library on the system where it is run.
- **Dynamic** linking is accomplished by placing the name of a sharable library in the executable image.
 - Actual linking with the library routines occurs when the image is run when executable and library are placed in memory.
 - An advantage of dynamic linking is that multiple programs can share a single copy of the library.

Static and dynamic linking with gcc

- By default executables are dynamically linked:

```
gcc app.c -o app
```

- To link statically, specify the `-static` argument to gcc:

```
gcc -static app.c -o app
```

- The `file` command is useful to determine the type of linking:

```
$ file app1
test: ELF 64-bit LSB executable, x86-64, version 1 (SYSV),
dynamically linked, ...

$ file app2
app2: ELF 64-bit LSB executable, x86-64, version 1 (GNU/Linux),
statically linked, ...
```

References

- Executable and Linkable Format
 - https://en.wikipedia.org/wiki/Executable_and_Linkable_Format
- Linker scripts
 - http://wiki.osdev.org/Linker_Scripts
 - https://www.math.utah.edu/docs/info/ld_3.html
- Linux ELF Object File Format Basics
 - <http://www.thegeekstuff.com/2012/07/elf-object-file-format>
- The ELF Object File Format by Dissection
 - <http://www.linuxjournal.com/article/1060>